

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
from datetime import timedelta
import datetime as dt
import matplotlib.pyplot as plt
import xgboost as xgb
from sklearn.cluster import MiniBatchKMeans
import warnings
from sklearn.model_selection import train_test_split
```

```
In [2]: data = pd.read_csv('green_tripdata_2016-12.csv', index_col=False)
data
```

Out[2]:

	VendorID	lpep_pickup_datetime	lpep_dropoff_datetime	store_and_fwd_flag	RatecodeID	PULocationID	DOLocationID	passenger_
0	2	2016-12-01 00:00:54	2016-12-01 00:06:54	N	1	92	192	
1	2	2016-12-01 00:52:41	2016-12-01 00:54:51	N	1	92	171	
2	2	2016-12-01 00:10:39	2016-12-01 00:14:47	N	1	75	238	
3	2	2016-12-01 00:12:16	2016-12-01 00:15:31	N	1	166	151	
4	2	2016-12-01 00:29:22	2016-12-01 00:39:51	N	1	166	42	
...
1224153	1	2016-12-31 23:00:16	2016-12-31 23:05:30	N	1	74	75	
1224154	1	2016-12-31 23:00:20	2016-12-31 23:04:05	N	1	42	41	
1224155	1	2016-12-31 23:00:08	2016-12-31 23:15:57	N	1	243	159	
1224156	1	2016-12-31 23:00:00	2016-12-31 23:10:15	N	1	244	120	
1224157	1	2016-12-31 23:00:10	2016-12-31 23:04:17	N	1	7	7	

1224158 rows × 9 columns

- VendorID - A code indicating the LPEP provider that provided the record. (1= Creative Mobile Technologies, LLC; 2= VeriFone Inc).
- lpep_pickup_datetime - The date and time when the meter was engaged.
- lpep_dropoff_datetime - The date and time when the meter was disengaged.

- `store_and_fwd_flag` - This flag indicates whether the trip record was held in vehicle memory before sending to the vendor, aka “store and forward,” because the vehicle did not have a connection to the server. ('Y' = store and forward trip 'N' = not a store and forward trip)
- `RatecodeID` - The final rate code in effect at the end of the trip. (1= Standard rate 2=JFK 3=Newark 4=Nassau or Westchester 5=Negotiated fare 6=Group ride)
- `PULocationID` - pickup location matching zone numbers to the map
- `DOLocationID` - drop location matching zone numbers to the map
- `passenger_count` - The number of passengers in the vehicle. This is a driver-entered value.
- `trip_distance` - The elapsed trip distance in miles reported by the taximeter.
- `fare_amount` - The time-and-distance fare calculated by the meter.
- `extra` - Miscellaneous extras and surcharges.
- `mta_tax` - \$0.50 MTA tax that is automatically triggered based on the metered rate in use.
- `tip_amount` - **This field is automatically populated for credit card tips. Cash tips are not included.**
- `tolls_amount` - Total amount of all tolls paid in trip.
- `ehail_fee` - None
- `improvement_surcharge` - \$0.30 improvement surcharge assessed on hailed trips at the flag drop. The improvement surcharge began being levied in 2015.
- `total_amount` - The total amount charged to passengers. Does not include cash tips.
- `payment_type` - A numeric code signifying how the passenger paid for the trip. (1= Credit card 2= Cash 3= No charge 4= Dispute 5= Unknown 6= Voided trip)
- `trip_type` - A code indicating whether the trip was a street-hail or a dispatch that is automatically assigned based on the metered rate in use but can be altered by the driver. (1= Street-hail 2= Dispatch)

1. Characterize the data and comment about its quality

```
In [3]: data.shape
```

```
Out[3]: (1224158, 19)
```

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1224158 entries, 0 to 1224157
Data columns (total 19 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   VendorID              1224158 non-null  int64
 1   lpep_pickup_datetime  1224158 non-null  object
 2   lpep_dropoff_datetime 1224158 non-null  object
 3   store_and_fwd_flag    1224158 non-null  object
 4   RatecodeID            1224158 non-null  int64
 5   PULocationID          1224158 non-null  int64
 6   DOLocationID          1224158 non-null  int64
 7   passenger_count       1224158 non-null  int64
 8   trip_distance         1224158 non-null  float64
 9   fare_amount           1224158 non-null  float64
10   extra                 1224158 non-null  float64
11   mta_tax               1224158 non-null  float64
12   tip_amount            1224158 non-null  float64
13   tolls_amount          1224158 non-null  float64
14   ehail_fee             0 non-null        float64
15   improvement_surcharge 1224158 non-null  float64
16   total_amount          1224158 non-null  float64
17   payment_type          1224158 non-null  int64
18   trip_type             1224158 non-null  int64
dtypes: float64(9), int64(7), object(3)
memory usage: 177.5+ MB
```

```
In [5]: data.describe()
```

```
Out[5]:
```


	VendorID	RatecodeID	PULocationID	DOLocationID	passenger_count	trip_distance	fare_amount	extra	mta_
count	1.224158e+06	1.224158e+06	1.224158e+06	1.224158e+06	1.224158e+06	1.224158e+06	1.224158e+06	1.224158e+06	1.224158e
mean	1.799278e+00	1.085079e+00	1.131729e+02	1.290862e+02	1.362447e+00	2.618179e+00	1.159368e+01	3.629562e-01	4.872578e
std	4.005404e-01	5.660195e-01	7.628387e+01	7.723367e+01	1.036924e+00	2.806160e+00	9.793499e+00	3.901575e-01	8.517500e
min	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	0.000000e+00	0.000000e+00	-4.990000e+02	-4.500000e+00	-5.000000e
25%	2.000000e+00	1.000000e+00	4.900000e+01	6.100000e+01	1.000000e+00	9.800000e-01	6.000000e+00	0.000000e+00	5.000000e
50%	2.000000e+00	1.000000e+00	8.300000e+01	1.290000e+02	1.000000e+00	1.700000e+00	9.000000e+00	5.000000e-01	5.000000e
75%	2.000000e+00	1.000000e+00	1.730000e+02	1.930000e+02	1.000000e+00	3.200000e+00	1.400000e+01	5.000000e-01	5.000000e
max	2.000000e+00	6.000000e+00	2.650000e+02	2.650000e+02	9.000000e+00	1.475000e+02	1.007000e+03	4.500000e+00	5.000000e

- There're 1224158 data samples with 19 features. There're 3 objective columns, the others are continuous variables. 'ehail_fee' are all NAN, no missing values in this table other than that.
- There're negative values in 'fare_amount', 'extra', 'mta_tax', 'tip_amount', 'tolls_amount', 'improvement_surcharge' and 'total_amount', and value 0 in 'passenger_count', 'trip_distance', which makes no sense.
- The maximum trip_distance of 147.5 miles, tip_amount of 250.7 dollars, tolls_amount of 298 dollars, total_amount of 1008.3 dollars are also weird.

2. Visualize and Clean Data

Trip Distance

```
In [6]: plt.figure(figsize=(8,6))  
plt.scatter(range(data.shape[0]), np.sort(data["trip_distance"].values), alpha=0.2)  
plt.xlabel('index', fontsize=12)  
plt.ylabel('trip_distance', fontsize=12)  
plt.title('Trip Distance')  
plt.show()
```



```
In [7]: data["trip_distance"].quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```

```
Out[7]: 0.00      0.00
        0.01      0.00
        0.25      0.98
        0.50      1.70
        0.75      3.20
        0.99     13.97
        1.00     147.50
        Name: trip_distance, dtype: float64
```

The 99 quantile is 13.97 miles while the maximum is 147.5 miles


clipping trip distance longer than 20 miles and shorter than 0.1 mile.

```
In [8]: data = data.loc[(data["trip_distance"]<=20)&(data["trip_distance"]>0.1)].reset_index(drop=True)
```

Trip Duration

```
In [9]: data.lpep_pickup_datetime = pd.to_datetime(data.lpep_pickup_datetime)
        data.lpep_dropoff_datetime = pd.to_datetime(data.lpep_dropoff_datetime)
        data['duration'] = (data.lpep_dropoff_datetime - data.lpep_pickup_datetime).dt.total_seconds() / 60
```

```
In [10]: plt.figure(figsize=(8,6))
plt.scatter(range(data.shape[0]), np.sort(data["duration"].values), alpha=0.2)
plt.xlabel('index', fontsize=12)
plt.ylabel('duration', fontsize=12)
plt.title('Duartion')
plt.show()
```



```
In [11]: data["duration"].quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```

```
Out[11]: 0.00      0.000000
          0.01      1.583333
          0.25      5.883333
          0.50      9.750000
          0.75     16.116667
          0.99     64.433333
          1.00    2487.900000
          Name: duration, dtype: float64
```

The 99 quantile is 64.5 min while the maximum is 1439.9 min

clipping trip distance longer than 60 min and less that one minute


```
In [12]: data = data.loc[(data["duration"]<=60)&(data["duration"]>=1)].reset_index(drop=True)
```

Speed

```
In [13]: data['speed'] = data.trip_distance/data.duration
```



```
In [14]: plt.figure(figsize=(8,6))  
plt.scatter(range(data.shape[0]), np.sort(data["speed"].values), alpha=0.2)  
plt.xlabel('index', fontsize=12)  
plt.ylabel('speed', fontsize=12)  
plt.title('Speed')  
plt.show()
```



```
In [15]: data.speed.describe()
```


```
Out[15]: count      1.181390e+06  
         mean       2.057084e-01  
         std        1.003032e-01  
         min        2.071346e-03  
         25%        1.496104e-01  
         50%        1.876161e-01  
         75%        2.376238e-01  
         max        1.419048e+01  
         Name: speed, dtype: float64
```

Delete records with speed more than 1 mile/min.

```
In [16]: data = data.loc[data.speed<=1].reset_index(drop=True)
```

Passenger Count

```
In [17]: plt.hist(data.passenger_count)
plt.title('Passenger Count')
plt.xlabel('Passenger', fontsize=12)
plt.show()
```



```
In [18]: data["passenger_count"].value_counts()
```

```
Out[18]: 1    997993
2     89920
5     42155
3     23273
6     20086
4      7666
0         29
8          8
7          7
9          1
Name: passenger_count, dtype: int64
```


```
In [19]: data["passenger_count"].quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```

```
Out[19]: 0.00    0.0  
         0.01    1.0  
         0.25    1.0  
         0.50    1.0  
         0.75    1.0  
         0.99    6.0  
         1.00    9.0  
Name: passenger_count, dtype: float64
```

Delete records with 0 pssenger and passenger_count greater than 6.

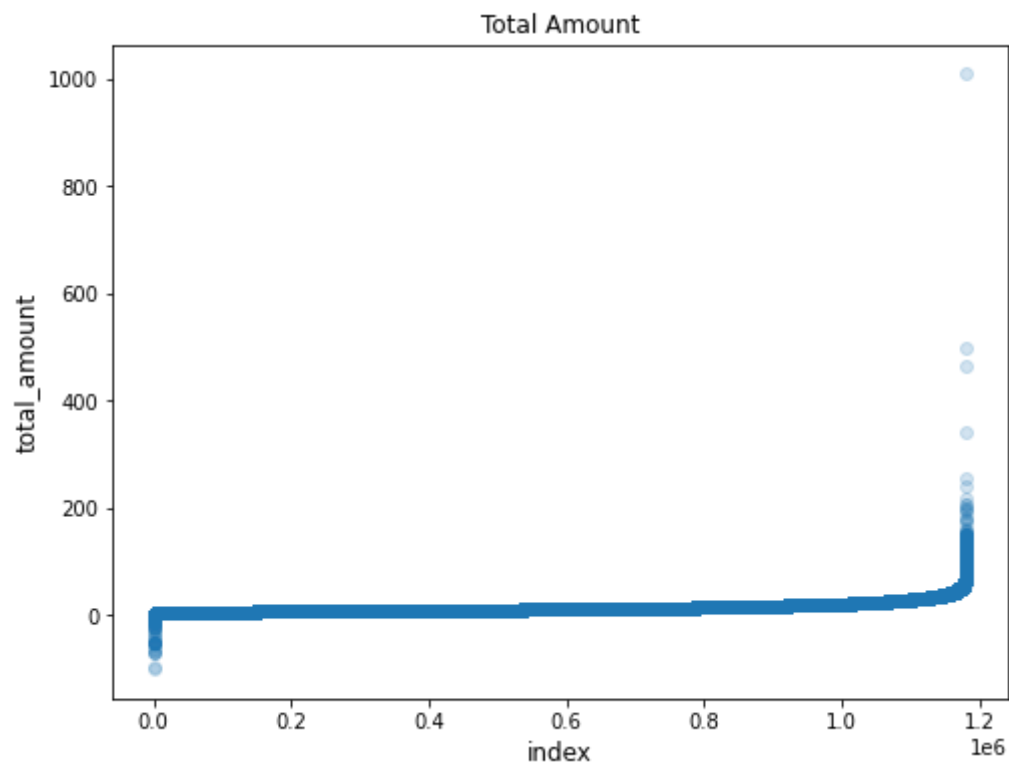
```
In [20]: data = data.loc[(data["passenger_count"]>0) & (data["passenger_count"]<=6)].reset_index(drop=True)
```

```
In [21]: plt.hist(data.passenger_count)  
plt.title('Passenger Count')  
plt.xlabel('Passenger', fontsize=12)  
plt.show()
```



Total Amount

```
In [22]: plt.figure(figsize=(8,6))
plt.scatter(range(data.shape[0]), np.sort(data["total_amount"].values), alpha=0.2)
plt.xlabel('index', fontsize=12)
plt.ylabel('total_amount', fontsize=12)
plt.title('Total Amount')
plt.show()
```



```
In [23]: data["total_amount"].quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```

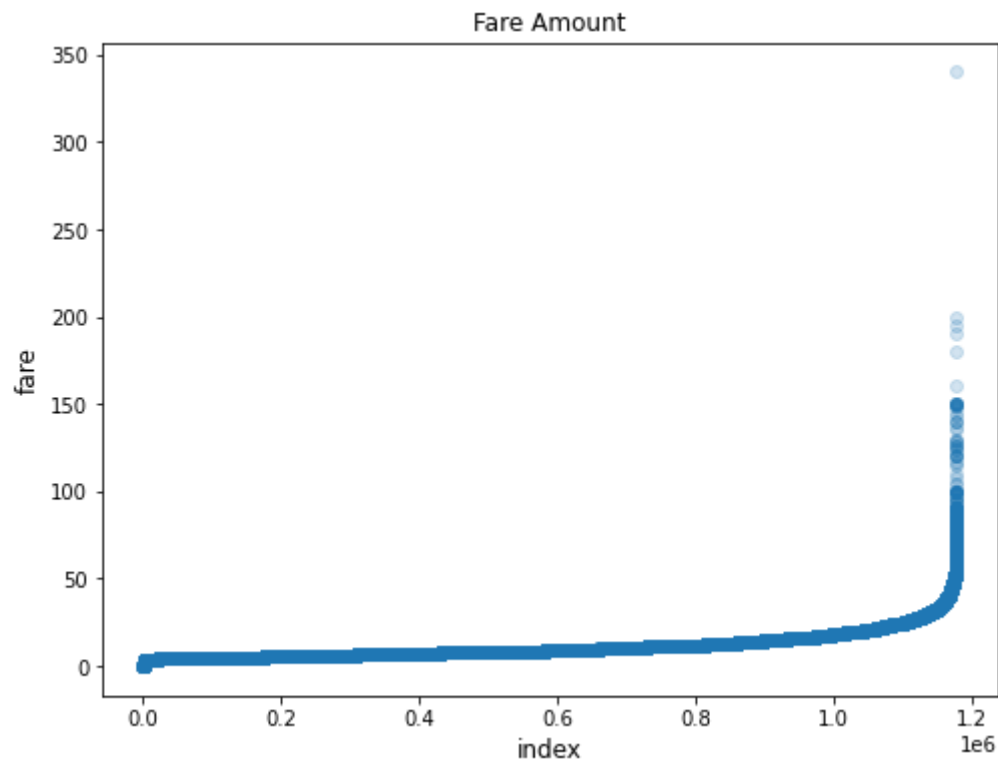
```
Out[23]: 0.00    -100.0
0.01      4.3
0.25      7.8
0.50     10.8
0.75     16.3
0.99     47.3
1.00    1008.3
Name: total_amount, dtype: float64
```

Delete total amount larger than 400 and less than 0

```
In [24]: data = data.loc[(data["total_amount"] <= 400) & (data["total_amount"] > 0)].reset_index(drop=True)
```

Fare Amount

```
In [25]: plt.figure(figsize=(8,6))
plt.scatter(range(data.shape[0]), np.sort(data["fare_amount"].values), alpha=0.2)
plt.xlabel('index', fontsize=12)
plt.ylabel('fare', fontsize=12)
plt.title('Fare Amount')
plt.show()
```



```
In [26]: data["fare_amount"].quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```


```
Out[26]: 0.00      0.0  
         0.01      3.5  
         0.25      6.0  
         0.50      9.0  
         0.75     14.0  
         0.99     40.0  
         1.00    340.0  
         Name: fare_amount, dtype: float64
```

Delete records with 0 fare amount.

```
In [27]: data = data.loc[data.fare_amount>0].reset_index(drop=True)
```

Tip Amount

```
In [28]: plt.figure(figsize=(8,6))
plt.scatter(range(data.shape[0]), np.sort(data["tip_amount"].values), alpha=0.2)
plt.xlabel('index', fontsize=12)
plt.ylabel('tip', fontsize=12)
plt.title('Tip Amount')
plt.show()
```



```
In [29]: data["tip_amount"].quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```

```
Out[29]: 0.00      0.00
0.01      0.00
0.25      0.00
0.50      0.00
0.75      1.86
0.99      7.84
1.00     240.00
Name: tip_amount, dtype: float64
```


Delete records with which tips is larger than fare amount

```
In [30]: data = data.loc[data.tip_amount < data.fare_amount].reset_index(drop=True)
```

```
In [31]: data.describe()
```

Out[31]:


	VendorID	RatecodeID	PULocationID	DOLocationID	passenger_count	trip_distance	fare_amount	extra	mta_t
count	1.176458e+06	1.176458e+06	1.176458e+06	1.176458e+06	1.176458e+06	1.176458e+06	1.176458e+06	1.176458e+06	1.176458e+06
mean	1.800384e+00	1.045240e+00	1.127070e+02	1.290701e+02	1.363125e+00	2.580974e+00	1.131171e+01	3.669749e-01	4.944996e-01
std	3.997120e-01	4.155648e-01	7.613657e+01	7.720240e+01	1.038978e+00	2.482266e+00	7.698733e+00	3.794237e-01	5.215316e-01
min	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.100000e-01	1.000000e-02	0.000000e+00	0.000000e+00
25%	2.000000e+00	1.000000e+00	4.900000e+01	6.100000e+01	1.000000e+00	1.000000e+00	6.000000e+00	0.000000e+00	5.000000e-01
50%	2.000000e+00	1.000000e+00	8.200000e+01	1.290000e+02	1.000000e+00	1.710000e+00	9.000000e+00	5.000000e-01	5.000000e-01
75%	2.000000e+00	1.000000e+00	1.680000e+02	1.930000e+02	1.000000e+00	3.200000e+00	1.400000e+01	5.000000e-01	5.000000e-01
max	2.000000e+00	6.000000e+00	2.650000e+02	2.650000e+02	6.000000e+00	2.000000e+01	3.400000e+02	4.500000e+00	5.000000e-01

3. Find interesting trip statistics grouped by hour

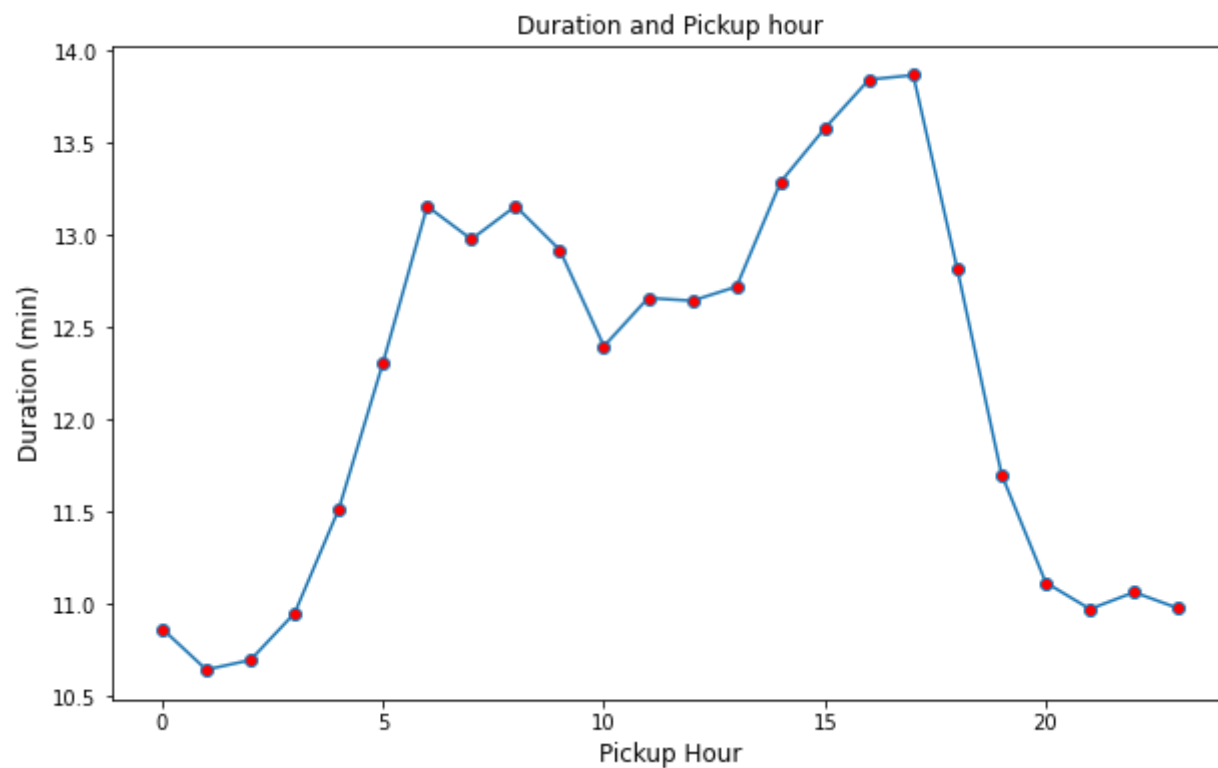
```
In [32]: data['pickup_hour'] = data['lpep_pickup_datetime'].dt.hour
data['dropoff_hour'] = data['lpep_dropoff_datetime'].dt.hour
data['pickup_weekday'] = data['lpep_pickup_datetime'].dt.weekday
data['dropoff_weekday'] = data['lpep_dropoff_datetime'].dt.weekday
```

```
In [33]: metric = 'pickup_hour'
y = 'duration'
data.boxplot(column=y, by=[metric], figsize=(8,5))
```

```
Out[33]: <AxesSubplot:title={'center':'duration'}, xlabel='[pickup_hour] '>
```



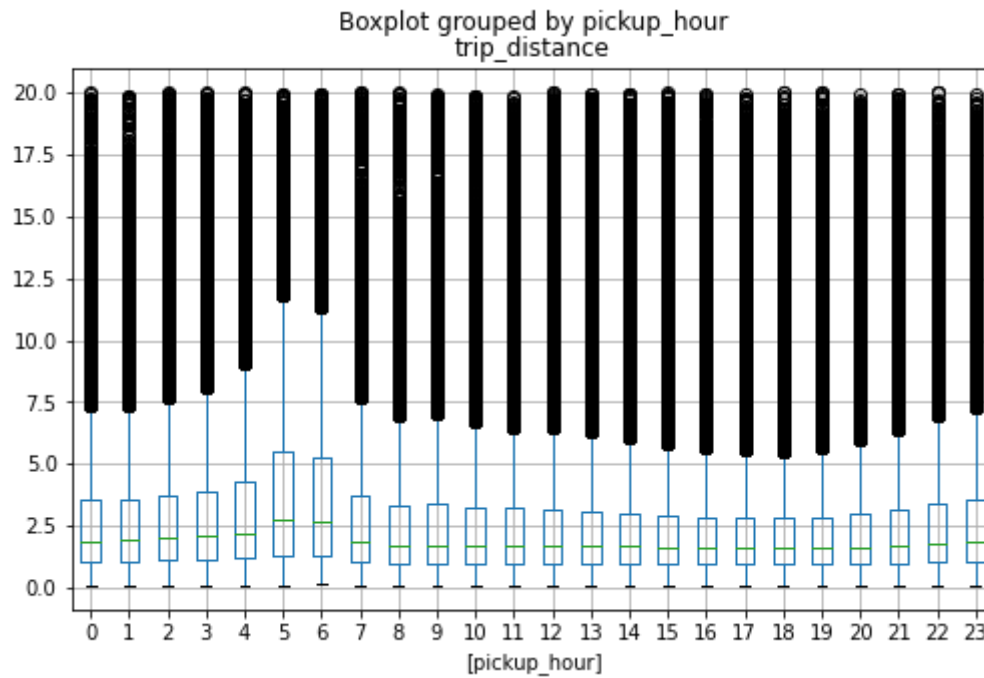
```
In [34]: duration_hour = data.groupby(by='pickup_hour',as_index=False)['duration'].mean()  
plt.figure(figsize=(10,6))  
plt.plot(duration_hour['pickup_hour'], duration_hour['duration'],marker='o', markerfacecolor='red',label=  
plt.xlabel('Pickup Hour', fontsize=12)  
plt.ylabel('Duration (min)', fontsize=12)  
plt.title('Duration and Pickup hour')  
plt.show()
```




The trip duration at night (7pm - 4am) is shorter than any other time during the day.

```
In [35]: metric = 'pickup_hour'
y = 'trip_distance'
data.boxplot(column=y, by=[metric], figsize=(8,5))
```

```
Out[35]: <AxesSubplot:title={'center':'trip_distance'}, xlabel='[pickup_hour] '>
```



```
In [36]: distance_hour = data.groupby(by='pickup_hour',as_index=False)['trip_distance'].mean()  
plt.figure(figsize=(10,6))  
plt.plot(distance_hour['pickup_hour'], distance_hour['trip_distance'], marker='o', markerfacecolor='red')  
plt.xlabel('Pickup Hour', fontsize=12)  
plt.ylabel('Distance (mile)', fontsize=12)  
plt.title('Distance and Pickup hour')  
plt.show()
```



Trip distance around 5am-6am is more likely to be longer than any other time during the day.

4. What kind of trip yields better tips?

According to the data dictionary on the website, the records only include tip information paid by credit card. So delete records with other payment method.

When defining better tips, it's more reasonable to evaluate tip percentage other than its absolute value. So the model is built to predict which trip yields a higher tip percentage.


```
In [37]: data['tip_percent'] = data.tip_amount/data.fare_amount
```

```
In [38]: df = data.loc[data.payment_type==1].reset_index(drop=True)
```

```
In [39]: df.tip_percent.quantile([.0, 0.01, 0.25, .5, 0.75, 0.99, 1])
```

```
Out[39]: 0.00    0.000000
          0.01    0.000000
          0.25    0.142857
          0.50    0.217931
          0.75    0.252000
          0.99    0.440741
          1.00    0.993333
          Name: tip_percent, dtype: float64
```

```
In [40]: plt.figure(figsize=(8,6))
plt.scatter(range(df.shape[0]), np.log(np.sort(df["tip_percent"].values)+1), alpha=0.2)
plt.xlabel('index', fontsize=12)
plt.ylabel('tip_percent', fontsize=12)
plt.show()
```



```
In [41]: loction = pd.read_csv('taxi+_zone_lookup.csv')
loction
```

Out[41]:

	LocationID	Borough	Zone	service_zone
0	1	EWR	Newark Airport	EWR
1	2	Queens	Jamaica Bay	Boro Zone
2	3	Bronx	Allerton/Pelham Gardens	Boro Zone
3	4	Manhattan	Alphabet City	Yellow Zone
4	5	Staten Island	Arden Heights	Boro Zone
...
260	261	Manhattan	World Trade Center	Yellow Zone
261	262	Manhattan	Yorkville East	Yellow Zone
262	263	Manhattan	Yorkville West	Yellow Zone
263	264	Unknown	NV	NaN
264	265	Unknown	NaN	NaN

265 rows × 4 columns

Map the LocationID to borough.

```
In [42]: d = {'Manhattan':1, 'Queens':2, 'Brooklyn':3, 'Bronx':4, 'Staten Island':5, 'Unknown':6, 'EWR':7}
df['PUL'] = df.PULocationID.map(lambda x:loction.Borough[x-1]).map(d)
df['DOL'] = df.DOLocationID.map(lambda x:loction.Borough[x-1]).map(d)
```

```
In [43]: df['store_and_fwd_flag'] = df['store_and_fwd_flag'].map({'N':0, 'Y':1})
```



```
In [44]: df.columns
```

```
Out[44]: Index(['VendorID', 'lpep_pickup_datetime', 'lpep_dropoff_datetime',  
              'store_and_fwd_flag', 'RatecodeID', 'PULocationID', 'DOLocationID',  
              'passenger_count', 'trip_distance', 'fare_amount', 'extra', 'mta_tax',  
              'tip_amount', 'tolls_amount', 'ehail_fee', 'improvement_surcharge',  
              'total_amount', 'payment_type', 'trip_type', 'duration', 'speed',  
              'pickup_hour', 'dropoff_hour', 'pickup_weekday', 'dropoff_weekday',  
              'tip_percent', 'PUL', 'DOL'],  
              dtype='object')
```

```
In [45]: ind = [0, 3, 4,15,16] + list(range(7, 13)) + list(range(18,22)) + [23,24,25,26,27]  
features = df.columns[ind]  
features
```

```
Out[45]: Index(['VendorID', 'store_and_fwd_flag', 'RatecodeID', 'improvement_surcharge',  
              'total_amount', 'passenger_count', 'trip_distance', 'fare_amount',  
              'extra', 'mta_tax', 'tip_amount', 'trip_type', 'duration', 'speed',  
              'pickup_hour', 'pickup_weekday', 'dropoff_weekday', 'tip_percent',  
              'PUL', 'DOL'],  
              dtype='object')
```

```
In [46]: train,val = train_test_split(df[features],test_size=0.2)
```

Target Encoding

```
In [47]: for col in ['pickup_weekday','pickup_hour','PUL', 'DOL']:  
        gby = train.groupby(col).mean()[['duration','tip_amount']]  
        gby.columns = ['%s_gby_%s' % (c, col) for c in gby.columns]  
        train = pd.merge(train, gby, how='left', left_on=col, right_index=True)  
        val = pd.merge(val, gby, how='left', left_on=col, right_index=True)
```

```
In [48]: X_train = train.drop(columns=['tip_amount','tip_percent'])  
y_train = train['tip_percent']  
X_val = val.drop(columns=['tip_amount','tip_percent'])  
y_val = val['tip_percent']
```

XGBoost Model

```
In [49]: xgb_pars = {'min_child_weight': 50, 'eta': 0.1, 'colsample_bytree': 0.3, 'max_depth': 10,
                    'subsample': 0.8, 'lambda': 1., 'nthread': -1, 'booster': 'gbtree',
                    'eval_metric': 'rmse', 'objective': 'reg:squarederror'}
```

```
In [50]: dtrain = xgb.DMatrix(X_train.values, label=y_train)
dvalid = xgb.DMatrix(X_val.values, label=y_val)
watchlist = [(dtrain, 'train'), (dvalid, 'valid')]
```

```
In [51]: t0 = dt.datetime.now()
gbm = xgb.train(xgb_pars, dtrain, 600, watchlist, early_stopping_rounds=50,
               maximize=False, verbose_eval=10)
t1 = dt.datetime.now()
print('Time fitting xgb: %i seconds' % (t1 - t0).seconds)
```

[1]	train-rmse:0.26846	valid-rmse:0.26866
[2]	train-rmse:0.24300	valid-rmse:0.24321
[3]	train-rmse:0.22269	valid-rmse:0.22293
[4]	train-rmse:0.20475	valid-rmse:0.20504
[5]	train-rmse:0.18837	valid-rmse:0.18870
[6]	train-rmse:0.17444	valid-rmse:0.17483
[7]	train-rmse:0.16228	valid-rmse:0.16272
[8]	train-rmse:0.15169	valid-rmse:0.15219
[9]	train-rmse:0.14254	valid-rmse:0.14310
[11]	train-rmse:0.12579	valid-rmse:0.12644
[12]	train-rmse:0.11993	valid-rmse:0.12064
[13]	train-rmse:0.11497	valid-rmse:0.11575
[14]	train-rmse:0.11085	valid-rmse:0.11167
[15]	train-rmse:0.10734	valid-rmse:0.10820
[16]	train-rmse:0.10440	valid-rmse:0.10530
[17]	train-rmse:0.10193	valid-rmse:0.10289
[18]	train-rmse:0.09985	valid-rmse:0.10088
[19]	train-rmse:0.09413	valid-rmse:0.09512
[21]	train-rmse:0.09149	valid-rmse:0.09258
[22]	train-rmse:0.08840	valid-rmse:0.08861

The final RMSE for the validation test to predict a high tip percentage reaches 0.027.

Feature Importance

```
In [52]: features = X_train.columns
```

```
In [53]: feature_importance_dict = gbm.get_fscore()
fs = ['f%i' % i for i in range(len(features))]
f1 = pd.DataFrame({'f': list(feature_importance_dict.keys()), 'importance': list(feature_importance_dict.values())})
f2 = pd.DataFrame({'f': fs, 'feature_name': features})
feature_importance = pd.merge(f1, f2, how='right', on='f')
feature_importance = feature_importance.fillna(0)
```

```
In [54]: feature_importance.sort_values('importance', ascending=False)
```

Out[54]:

	f	importance	feature_name
4	f4	20864	total_amount
11	f11	19181	duration
6	f6	17319	trip_distance
12	f12	14991	speed
7	f7	11606	fare_amount
13	f13	9674	pickup_hour
21	f21	8578	tip_amount_gby_pickup_hour
20	f20	6901	duration_gby_pickup_hour
14	f14	6691	pickup_weekday
15	f15	6084	dropoff_weekday
19	f19	4495	tip_amount_gby_pickup_weekday
24	f24	4145	duration_gby_DOL
18	f18	4104	duration_gby_pickup_weekday
17	f17	3919	DOL
5	f5	3784	passenger_count
16	f16	3534	PUL
22	f22	3152	duration_gby_PUL
23	f23	2895	tip_amount_gby_PUL
25	f25	2587	tip_amount_gby_DOL
0	f0	2373	VendorID
8	f8	1881	extra
2	f2	723	RatecodeID
9	f9	384	mta_tax
10	f10	358	trip_type

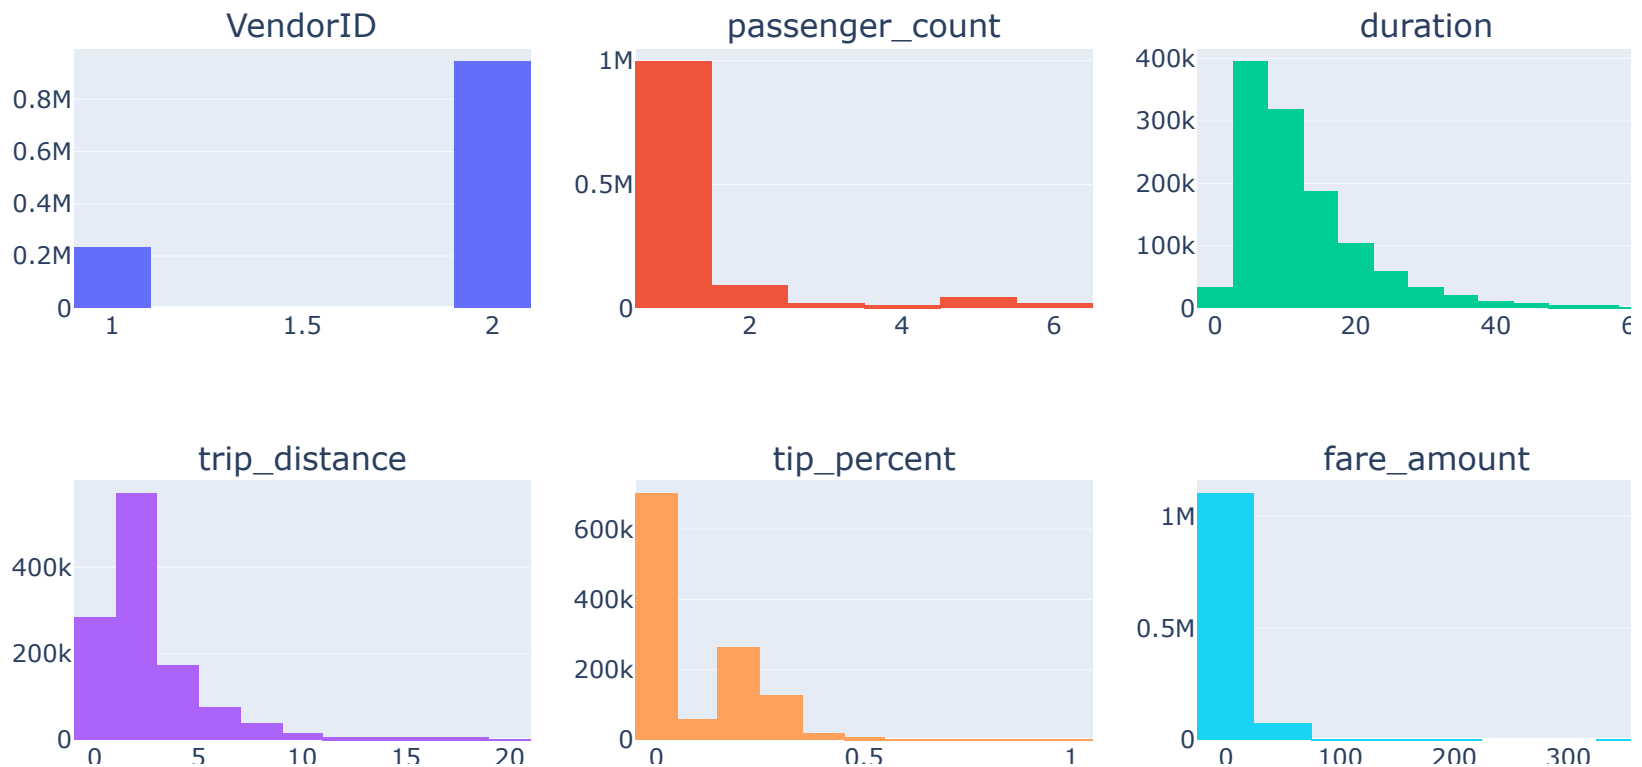
	f	importance	feature_name
3	f3	345	improvement_surcharge
1	f1	83	store_and_fwd_flag

5. Visualize the data to understand trip patterns

```
In [55]: import plotly.graph_objects as go
from plotly.subplots import make_subplots

fig = make_subplots(rows=2,
                    cols=3,
                    subplot_titles=['VendorID', 'passenger_count', 'duration', 'trip_distance', 'tip_percent', 'fare_amount'])

fig.append_trace(go.Histogram(x=data['VendorID'], nbinsx=10), 1, 1)
fig.append_trace(go.Histogram(x=data['passenger_count'], nbinsx=10), 1, 2)
fig.append_trace(go.Histogram(x=data['duration'], nbinsx=20), 1, 3)
fig.append_trace(go.Histogram(x=data['trip_distance'], nbinsx=10), 2, 1)
fig.append_trace(go.Histogram(x=data['tip_percent'], nbinsx=10), 2, 2)
fig.append_trace(go.Histogram(x=data['fare_amount'], nbinsx=10), 2, 3)
fig.show()
```



- **Among the cleaned data, VeriFone Inc provides about 4 times records than Creative Mobile Technologies, LLC.**
- **1 passenger is the most common scenario.**
- **Most trips last around 2.5 min to 20 min.**
- **Most trip distances are around 1 mile to 3 miles.**
- **Most passengers won't pay the tip. Among tip payers, people are most willingly to pay 15% - 25% fare amount fee as the tip.**
- **Most fare amount are under \$25.**